

Experiment 1 : Implementation of Linear and Logistic Regression on a Real-World Dataset

```
# =====
# STEP 1: Import Required Libraries
# =====
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression, LogisticRegression
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.metrics import confusion_matrix, accuracy_score

# =====
# STEP 2: Load Dataset
# =====
df = pd.read_csv('/mnt/data/student_performance (1).csv')
print(df.head())

# =====
# STEP 3: LINEAR REGRESSION
# Predict Final_Score
# =====
X = df[['Hours_Studied', 'Attendance', 'Assignment_Score', 'Midterm_Score']]
y = df['Final_Score']

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)

lr = LinearRegression()
lr.fit(X_train, y_train)

y_pred = lr.predict(X_test)

# -----
# Graph 1: Actual vs Predicted
# -----
plt.figure()
plt.scatter(y_test, y_pred)
plt.xlabel("Actual Final Score")
plt.ylabel("Predicted Final Score")
```

```

plt.title("Linear Regression: Actual vs Predicted")
plt.show()

# -----
# Graph 2: Residual Plot
# -----
residuals = y_test - y_pred

plt.figure()
plt.scatter(y_pred, residuals)
plt.axhline(0)
plt.xlabel("Predicted Final Score")
plt.ylabel("Residuals")
plt.title("Residual Plot")
plt.show()

# =====
# STEP 4: LOGISTIC REGRESSION
# Pass / Fail Classification
# =====
df['Pass'] = (df['Final_Score'] >= 60).astype(int)

X = df[['Hours_Studied', 'Attendance', 'Assignment_Score', 'Midterm_Score']]
y = df['Pass']

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)

log_reg = LogisticRegression(max_iter=1000)
log_reg.fit(X_train, y_train)

y_pred = log_reg.predict(X_test)

# -----
# Graph 3: Confusion Matrix
# -----
cm = confusion_matrix(y_test, y_pred)

plt.figure()
sns.heatmap(cm, annot=True, fmt='d')
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Logistic Regression Confusion Matrix")

```

```

plt.show()

# -----
# Graph 4: Probability Distribution
# -----
y_prob = log_reg.predict_proba(X_test)[:, 1]

plt.figure()
plt.hist(y_prob, bins=10)
plt.xlabel("Probability of Passing")
plt.ylabel("Count")
plt.title("Pass Probability Distribution")
plt.show()

# =====
# STEP 5: Accuracy
# =====
accuracy = accuracy_score(y_test, y_pred)
print("Logistic Regression Accuracy:", accuracy)

```

'OUTPUT :

The screenshot shows a Jupyter Notebook interface with several code cells and their corresponding outputs.

- Cell 1:** Shows the import statements for pandas, matplotlib.pyplot, and seaborn.
- Cell 2:** Shows the import statement for sklearn.model_selection.
- Cell 3:** Shows the import statement for sklearn.linear_model.
- Cell 4:** Shows the import statements for mean_squared_error and r2_score from sklearn.metrics.
- Cell 5:** Shows the import statement for confusion_matrix and accuracy_score from sklearn.metrics.
- Cell 6:** Shows the code to read the 'StudentPerformance.csv' file and print its head.

The output of Cell 6 is a table showing the first 5 rows of the dataset:

	Hours Studied	Previous Scores	Extracurricular Activities	Sleep Hours
0	7	99	Yes	9
1	4	82	No	4
2	8	51	Yes	7
3	5	52	Yes	5
4	7	75	No	8


Below the table, there is another section of code and output:

	Sample Question Papers Practiced	Performance Index
0	1	91.0
1	2	65.0
2	2	45.0
3	2	36.0
4	5	66.0


```
[11] X = df[['Hours Studied', 'Previous Scores', 'Sleep Hours', 'Sample Question Papers Practiced']]  
y = df['Performance Index']  
  
X_train, X_test, y_train, y_test = train_test_split(  
    X, y, test_size=0.2, random_state=42  
)  
  
lr = LinearRegression()  
lr.fit(X_train, y_train)  
  
y_pred = lr.predict(X_test)
```

```
[12] nlt finnme/
```

```
[12] Os  
❶ plt.figure()  
plt.scatter(y_test, y_pred)  
plt.xlabel("Actual Final Score")  
plt.ylabel("Predicted Final Score")  
plt.title("Linear Regression: Actual vs Predicted")  
plt.show()
```




?



```
[13] ✓ 0s
    residuals = y_test - y_pred

    plt.figure()
    plt.scatter(y_pred, residuals)
    plt.axhline(0)
    plt.xlabel("Predicted Final Score")
    plt.ylabel("Residuals")
    plt.title("Residual Plot")
    plt.show()
```



```
[15] ✓ 0s
    df['Pass'] = (df['Performance Index'] >= 60).astype(int)

    X = df[['Hours Studied', 'Previous Scores', 'Sleep Hours', 'Sample Question Papers Practiced']]
    y = df['Pass']


    X_train, X_test, y_train, y_test = train_test_split(
        X, y, test_size=0.2, random_state=42
    )

    log_reg = LogisticRegression(max_iter=1000)
    log_reg.fit(X_train, y_train)

    y_pred = log_reg.predict(X_test)
```

```
[16] ✓ 0s
    cm = confusion_matrix(y_test, y_pred)


    plt.figure()
    sns.heatmap(cm, annot=True, fmt='d')
    plt.xlabel("Predicted")
    plt.ylabel("Actual")
    plt.title("Logistic Regression Confusion Matrix")
    plt.show()
```



```
[17] ✓ 0s
    y_prob = log_reg.predict_proba(X_test)[:, 1]
```

```
[17] y_prob = log_reg.predict_proba(X_test)[:, 1]

plt.figure()
plt.hist(y_prob, bins=10)
plt.xlabel("Probability of Passing")
plt.ylabel("Count")
plt.title("Pass Probability Distribution")
plt.show()
```



```
accuracy = accuracy_score(y_test, y_pred)
print("Logistic Regression Accuracy:", accuracy)
```